

CONTENTS

ABSTRACT	1
1. INTRODUCTION	1
2. INDUSTRIAL SETTING	2
2.1. SETTING FOUNDATIONS	3
2.2. TOWARDS A REASONING FRAMEWORK FOR	4
3. VADALOG REASONING	5
4. THE VADA-SA FRAMEWORK	5
4.1. THE ANONYMIZATION CYCLE AND THE	5
4.2. STATISTICAL DISCLOSURE RISK ESTIMATION	6
4.3. SMART ANONYMIZATION	8
4.4. ENHANCING ANONYMIZATION	9
5. EXPERIMENTS	9
5.1. TESTING ANONYMIZATION CAPABILITY	10
5.2. TESTING SCALABILITY	10
6. RELATED WORK	11
7. CONCLUSION	12
REFERENCES	12

Financial Data Exchange with Statistical Confidentiality: A Reasoning-based Approach*

Luigi Bellomarini
Banca d'Italia

Rosario Laurendi
Banca d'Italia

Livia Blasi
Banca d'Italia

Emanuel Sallinger
TU Wien and University of Oxford

ABSTRACT

Confidentiality is a crucial requirement in financial data exchange processes. On the one hand, rich microdata is needed for most AI applications, including banking supervision, anti-money laundering, etc. On the other hand, organizations may not be legally authorized to see particular data, e.g., personal data. Striking the right balance provides a number of challenges.

Motivated by our experience with the Central Bank of Italy, in this work we present Vada-SA, a reasoning-based framework for financial data exchange with statistical confidentiality. We present a production-ready and fully engineered framework, adopting a reasoning approach. The framework includes explicit consideration of the reasoning process, the business context and declarative transparency that puts the user in control. We show and discuss a number of risk measures and anonymization criteria, implemented and operated in practice.

1 INTRODUCTION

Confidentiality in financial data exchange has multiple facets and touches different business segments of the FinTech area. In *open banking* settings, where the increasingly frequent interactions between financial intermediaries motivated by the unbundling and rebundling of the banking process sees the interplay of many actors, each interested in utilizing the data about a specific portion of the process, but with limited or no access-rights to the identity of the involved customers; in European-level *banking supervision*, where data exchange between the European Central Bank and the National Central Banks needs to reveal situations that are highly critical in terms of the “financial health” of the banks, while the identity of the involved customers tends to be irrelevant; in *anti-money laundering*, where most modern approaches pinpoint fraudulent or collusive cases by inspecting high-level features of the considered actors, without accessing their identity before any judicial or law-enforcement action authorizes it; in *statistical and economic research*, with the more and more common establishment of national “Research Data Centers”, data archives used by financial authorities that wish to share relevant financial data with universities and research institutes while keeping personal data reserved. Moreover, it goes without saying that the *GDPR regulation* makes the attention to the confidential transfer of personal data a central topic in Europe.

As a matter of fact, financial and statistical authorities and intermediaries look at solutions to share their own *microdata*, i.e., non-aggregated data at the finest level of granularity, while striking a good balance between their statistical relevance and the

need to eliminate any possible trace of personal identities. Many situations arise in the financial segment in which a counterparty must at the same time see parts of the data (to carry out a portion of the process) and must not see other parts which they are not legally authorized to see, e.g., personal data.

This paper is motivated by our experience with the Central Bank of Italy, which, in its capacity of national central bank, banking supervision and oversight authority and Financial Intelligence Unit for Italy, is touched by the problem of *confidential financial data exchange* in all its perspectives. In this work, we present VADA-SA, the joint effort of the Applied Research Team of the Bank of Italy, TU Wien and the University of Oxford towards a reasoning-based approach to the problem.

The desiderata. We start by laying out the main desiderata for a state-of-the-art financial data exchange solution with confidentiality: (i) It should be *context aware* and take into consideration the specific business domain and the characteristics of the involved entities and features to evaluate the risk of a breach of confidentiality; (ii) At the same time it should be *schema independent*, and operate regardless of the specific dataset structure; (iii) It should be *preemptive*, in the sense that it should be able to analyze a given dataset to be exchanged and provide a confidentiality score beforehand, so that analysts can evaluate the risk of sharing it; (iv) It should be *active*, in the sense that whenever the confidentiality score is over a certain threshold (e.g., statistically inferred or defined by the domain experts), the solution should be able to alter the data and *anonymize* them so that the threshold is respected; (v) It should embody a *statistics-preserving* anonymization logic, by removing the minimum amount of information needed to guarantee confidentiality, while preserving the statistical soundness and relevance of data; (vi) It should be *fully explainable*, meaning that the confidentiality score of a candidate dataset as well as the reasons for specific anonymization choices should be completely understandable to domain experts; plus it should have a transparent semantics of confidentiality; (vii) It should be *business friendly*, by being extensible, IT-independent and at business level, i.e., domain experts should operate autonomously in defining new scoring criteria as well as anonymization logic in a high-level non-technical language; (viii) It should be *scalable* and able to handle increasing data volumes.

Statistical Disclosure Control. The area of *Statistical Disclosure Control* [26, 35, 37] (SDC) represents a relevant yardstick for our work. The SDC approach concentrates on *re-identification*, i.e., the possibility for an attacker to cross-link information it rightfully retains, in particular, every single tuple of a legitimately owned database, with other data sources so as to find out the underlying identities (of the involved people, companies and stakeholders in general). SDC adopts quantitative indicators to take decisions on data sharing by evaluating the *risk of re-identification* and balancing it with the measure of the statistical

*The views and opinions expressed in this paper are those of the authors and do not necessarily reflect the official policy or position of Banca d'Italia.

relevance of the data, so as to minimize the risk while maximizing the *statistical utility*. SDC also studies solutions to transform, namely *anonymize*, the data to be shared, balancing confidentiality and statistical relevance. Commonly adopted techniques, featured by widespread tools such as *sdcMicro* [9], *μ -ARGUS* [27], and *ARX* [33], aim at removing potential identifiers (sometimes known as *quasi-identifiers*) of the disclosed tuples and include *value suppression*, *aggregation*, and *generalization*.

Unfortunately, the approaches and the tools offered so far by the SDC community do not fulfill the desiderata of a full-fledged solution needed by the processes of financial companies and organizations, like the Bank of Italy. First, to the best of our knowledge, all the existing SDC techniques are schema dependent and anonymization risk assessment and anonymization programs are tightly coupled to the dataset structure. Then, SDC techniques are only based on value statistics within the dataset to be anonymized and are not context-aware, while it is our experience that the risk of disclosure highly depends on the characteristics of the source and target databases [18] as well as the surrounding business information, e.g., availability of specific cross-linking data, even at tuple level. As a consequence, SDC techniques tend to fall short of accuracy in this respect. Although the anonymization techniques of SDC put into action interesting ideas and, in general, preserve statistical relevance of the datasets, to the best of our knowledge, all of them lack full explainability, unacceptable for financial organizations with strong accountability constraints. The lack of explainability prevents effective feedback-based adaptivity and the improvement of disclosure control proceeds by trial and error. Furthermore, all the existing tools tend to be not business friendly: they adopt a technology- and IT-dependent language (e.g., R libraries or Java), often lack clear semantics (typically only informally explained in the documentation), require adopters to have a technical background and are hardly extensible. Finally, such tools are data-scientist oriented libraries and, while showing good performance, do not have formal scalability guarantees.

Contribution. In this work we present VADA-SA, a *reasoning-based* framework for financial data exchange with statistical confidentiality. It is based on our long-term experience in developing AI-enhanced data-driven solutions revolving around *logic-based reasoning*. In particular, this work builds on the VADALOG System [6], a state-of-the-art *reasoning system* leveraging the VADALOG language, a member of the Datalog[±] family [12], exhibiting very good characteristics of scalability and expressive power. In particular, we contribute as follows.

- We present a production-ready and fully engineered framework, VADA-SA, for financial data exchange with confidential privacy, adopting a reasoning approach. The enterprise data to be shared, along with the metadata, are modeled as the *extensional component* of the reasoning process, whereas standard risk measures and anonymization methods are modeled as the *intensional component* of the process, i.e., a set of VADALOG rules. The activation of the rules upon the extensional component —i.e., the *reasoning process*— produces the *derived extensional component*, which is either a fully explained risk measure for a given dataset to be exchanged or its anonymized version.
- We show and discuss (and through VADA-SA ship off-the-shelf) a number of risk measures and anonymization criteria and illustrate how they can be managed in VADALOG.
- We suggest that the surrounding *business context* relevant for accurate risk measures is awarely modeled within the intensional component in terms of VADALOG rules, which are at the same

time *schema independent* w.r.t. the structure of the datasets. Although the framework targets financial data as a primary application, the techniques we present are general and can be applied in any context requiring statistical confidentiality.

- We envision that the SDC techniques can be used as a solid theoretical basis to craft a statistically preserving anonymization logic, yet, unlike existing approaches, we model in a purely *declarative* way in terms of VADALOG rules.
- We embrace a *user-delegation approach*, in the sense that by means of a semantically clear, fully declarative, non-technical and IT-independent language (i.e., characteristics that VADALOG embodies by design [6]), we delegate specific users to writing their own criteria and encoding the business knowledge, with cost and operational savings.
- In our framework, we inherit a set of benefits from logic-based reasoning. In particular, we refer to the pros of declarative approaches that, unlike procedural programming, relieve the users from the need to understand the internals of anonymization methods when adopting it. Full explainability is guaranteed by standard logic entailment semantics, enforced with CHASE-based procedures [20] embodied in VADA-SA. Finally, the ideal balance between computational complexity and expressive power inherited from VADALOG, allows VADA-SA to achieve very good scalability.
- We discuss an interesting set of real-world risk measures and anonymization criteria, implemented and operated in practice.

Overview. The remainder of the paper is organized as follows. In Section 2 we pursue the industrial setting at the Bank of Italy. In Section 3 we introduce the background about VADALOG. Section 4 presents the VADA-SA framework and Section 5 shows it in action in relevant cases from the Bank of Italy. In Section 6 we discuss some related work and Section 7 concludes the paper.

2 INDUSTRIAL SETTING

The Bank of Italy has recently set up a *Research Data Center* (RDC).¹ At its core, there are a set of relational databases that store the *microdata*, i.e., the operational finest-grained data, from many core business applications such as the credit risk register, payment systems, balance of payments, banking supervision indicators, etc. The ultimate goal of RDC is sharing statistically relevant information with other cooperating institutions such as the National Statistical Office, other central banks, the European Central Bank, universities and research centers. While all these counter-parties operate within a “circle of trust”, and can thus access the mentioned microdata, the identities of the involved entities, be they companies, banks or people, should remain of the sole responsibility (and therefore visibility) of the Bank of Italy, which is legally in charge of the respective processing.

The microdata that the RDC deals with regard different business processes and originate from multiple sources, usually external to the Bank of Italy. These data are collected with a variety of methods such as statistical surveys or data flows and are organized into several *microdata DBs*, by business domain. The RDC aims at including 65 microdata DBs with operational data from 1977 to 2020 and expected size of 30-50TB, with a 1TB/month growth. The RDC currently stores 14 microdata DBs, about families and individuals, firms, and historical data, including:

- Household income and wealth
- Household finance and consumption

¹The RDC is part of the INEXDA initiative (<http://www.inexda.org/>) for the exchange of granular statistical data.

	I	Q	Q	Q	Q	Q	A	A	W
	Id	Area	Sector	Employees	Residential Rev.	Export Rev.	Exp. to DE	Grwth 6mos	W
1	612276	North	Public Service	50-200	0-30	0-30	30-60	2	230
2	737536	South	Commerce	201-1000	0-30	90+	0-30	-1	190
3	971906	Center	Commerce	1000+	0-30	30-60	0-30	4	70
4	589681	North	Textiles	1000+	90+	0-30	0-30	30	60
5	419410	North	Construction	1000+	90+	0-30	0-30	300	50
6	972915	North	Other	1000+	0-30	0-30	30-60	50	70
7	501118	North	Other	201-1000	60-90	90+	90+	-20	300
8	815363	North	Textiles	201-1000	60-90	30-60	90+	2	230
9	490065	South	Public Service	50-200	0-30	0-30	0-30	12	123
10	415487	South	Commerce	1000+	0-30	0-30	90+	3	145
11	399087	South	Commerce	50-200	30-60	0-30	30-60	2	70
12	170034	Center	Commerce	1000+	60-90	0-30	0-30	45	90
13	724905	Center	Construction	201-1000	0-30	30-60	0-30	2	200
14	554475	Center	Other	50-200	0-30	90+	0-30	0	104
15	946251	Center	Public Service	201-1000	30-60	90+	90+	150	30
16	581077	North	Textiles	50-200	0-30	60-90	30-60	-20	160
17	765562	South	Textiles	50-200	0-30	60-90	0-30	-7	200
18	154840	Center	Commerce	201-1000	0-30	60-90	0-30	4	220
19	600837	Center	Construction	50-200	0-30	60-90	0-30	20	190
20	220712	Center	Financial	1000+	30-60	60-90	30-60	-30	90

Figure 1: Microdata DB about inflation and growth.

- Financial literacy data
- Business outlook of industrial and service firms
- Italian housing market
- Inflation growth and expectations
- Historical archive of Italian credit.

Microdata DBs contain business data, including attributes that may disclose, directly or indirectly, the identity of the involved subjects; let us call these subjects *respondents*, by some abuse of the terminology adopted for statistical surveys. The risk for a tuple of a microdata DB to be associated (i.e., “linked”) to the respective real-world identity of the respondent is named *risk of re-identification*. Indeed, the notion of re-identification revolves around the (realistic) assumption that an external data source containing all the identities of the respondents exists; let us call *identity oracle* such database. The challenge here consists in mitigating the risk that an attacker could be able to link the value of some attributes of a tuple of the microdata DB, with those of a single tuple (or a very small set thereof) of the identity oracle and therefore disclose the respondent’s identity.

2.1 Setting Foundations

Let us frame our industrial context with the needed foundations.

Relational Foundations. Let C , N , and V be disjoint countably infinite sets of *constants*, (*labelled*) *nulls* and (regular) *variables*, respectively. A (relational) schema S is a finite set of relation symbols (or predicates) with associated arity. A *term* is either a constant or variable. An *atom* over S is an expression of the form $R(\bar{v})$, where $R \in S$ is of arity $n > 0$ and \bar{v} is an n -tuple of terms. A *database instance* (or simply *database*) over S associates to each relation symbol in S a relation of the respective arity over the domain of constants and nulls. The members of relations are called *tuples*. By some abuse of notations, we sometimes use the terms tuple and fact interchangeably.

The Microdata DB and the Identity Oracle. A microdata DB is a relation of schema $M(\bar{i}, \bar{q}, \bar{a}, W)$, where \bar{i} is an n -tuple of attributes defined as *direct identifiers*, \bar{q} is an n -uple of *quasi-identifiers*, \bar{a} is a set of *non-identifying* attributes and W is a *sampling weight*. An identity oracle is a relation of schema $O(\bar{i}', \bar{q}', I)$,

where \bar{i}' is a set of direct identifiers, \bar{q}' is a set of quasi-identifiers and I is the *identity* of the respondent.

- *Direct identifiers* are attributes s.t. their values (of each single attribute, separately) allow to determine the identity of the respondent, that is, for a given tuple of M , the join between M and O on an attribute of \bar{i} equated to an attribute of \bar{i}' selects a single tuple from O and therefore the resulting tuple discloses the respondent’s identity I . Observe that a direct identifier is a key attribute for O and it is assumed that $\bar{i} \subseteq \bar{i}'$. Examples of direct identifiers are the social security number, the Italian fiscal code, the driving licence number, etc.
- *Quasi-identifiers* are attributes s.t. the values of two or more of them, jointly, are likely to disclose the identity of the respondent, that is, for a given tuple of M , the join between M and O on two or more attributes of \bar{q} equated to attributes of \bar{q}' selects a small set of tuples of O and therefore likely discloses the respondent’s identity I . In other terms, quasi-identifiers are features that in specific combinations are enough selective to endanger the respondent’s confidentiality. This selectivity depends on the attribute (as some are intrinsically more specific) and, of course, on the combination of values, which can be more or less specific for a given context. For example, the joint use of age and address can be quite selective if we refer to a context of small dwellings, whereas gender and address would be less selective. On the other hand, occupation-gender is in general not very selective, whereas it can be extremely discriminating if we are referring to a context of a survey about gendered jobs in some country.
- *Non-identifying attributes* are those that do not fall in the two previous categories. These attributes are not critical because neither individually nor in combination with others, allow to disclose the identity of the respondent, i.e., re-identification is not possible. On the one hand, this can depend on an intrinsic scarce selectivity of the attribute like in the case of age in a given context, on the other hand, a non-identifying attribute can be even intrinsically identifying, yet its value certainly unknown to the identity oracle. This is the case, for instance, of internal system identifiers which are useless for re-identification.

- *Context and sampling weight.* We have touched on the notion of context when discussing quasi-identifiers, which are more or less selective depending on the domain of discourse. The context can be seen as a selection of tuples from O based on the domain of interest. For instance, if we were surveying the population of Milan, the only tuples of O referring to people living in Milan could be used to attempt re-identification of tuples of M , thus making it easier. The *sampling weight* accounts for the context by measuring the *representativeness* of a tuple t of M w.r.t. the entire context \mathbb{C} to which M refers. In this sense, R is a sample from O , and W_t is the tuple sampling weight.

There are different options for defining the sampling weight [7, 22]. The one we take inspiration from is the expected value of the number of entities having the same characteristics as t (according to a similarity function ϕ) in the sample distribution of O according to a given context \mathbb{C} . Given M , the weight W_t can be estimated for each tuple t from the posterior distribution of values for \bar{q} among the tuples. Many options are also possible for ϕ and the simplest one just uses equality of quasi-identifiers attributes. Higher weights denote statistically relevant tuples, likely carrying scarcely selective attributes; lower weights denote statistically less relevant tuples (outliers, as a limit case), likely with highly selective attributes.

- *Identity.* The value of such attribute stands for some universally recognized representation of one respondent's identity.

In our experience with the Bank of Italy, the categorization of microdata DB attributes as direct, quasi- and non-identifiers as well as weight estimation is a hybrid process involving human experience-based evaluation, learning from training sets, and domain-based reasoning, as we shall see.

2.2 Towards a Reasoning Framework for Statistical Disclosure Control

With the depicted context, we can achieve a straightforward definition of *re-identification risk* as the probability $\rho_t = 1/W_t$ of re-identifying t given the value of all its quasi-identifiers \bar{q} . We can say that, in some sense, provided that O is an abstraction, the sampling weight W_t is an estimator for the cardinality of the join $|\sigma_t(M) \bowtie_{\bar{q}} O|$, where σ denotes the selection and \bowtie the join.

However, re-identification risk is an upper bound for the real disclosure risk, in the assumption that all quasi-identifiers are known to a potential attacker. As a matter of fact, for a given tuple we may be interested in evaluating the risk only wrt a subset $\hat{q} \subset \bar{q}$ of quasi-identifiers, the ones we suppose the attacker is aware of or are more selective. Moreover, we may want to apply an arbitrary *risk weight* function λ , which takes as input W_t as well as the values for quasi-identifiers of t . Whence, the following definition of a general *statistical disclosure risk*:

$$\rho_{\hat{q}} = 1/\lambda(\sigma_{\hat{q}=\hat{q}}M) \quad (1)$$

The function λ computes an aggregate weight over the tuples selected by \hat{q} and generalizes many different risk measurement techniques, as we shall see, including the re-identification-based risk (for which $\lambda(\sigma_{\hat{q}=\hat{q}}R) = \sum_{t \in \sigma_{\hat{q}=\hat{q}}(R)} W_t$), but also *k-anonymity* [34], *individual risk* [7], and *SUDA* [19].

Use cases. Our industrial goal consists in: 1. evaluating the *statistical disclosure risk* and, 2. if unacceptable, take actions to counteract possible information disclosure, while preserving statistical significance (*anonymization*). The joint performance of the two mentioned actions is known as *statistical disclosure control* [18].

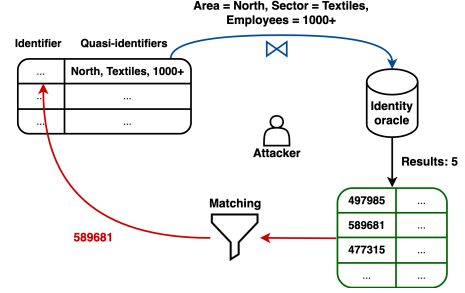


Figure 2: The attack strategy in action: by querying the identity oracle along Area, Sector and number of Employees, an attacker can narrow the search to few candidates and make a plausible guess about respondents' identity.

Figure 1 reports a fragment of a microdata DB of the Bank of Italy RDC, whose data derive from an *Inflation and Growth Survey*. This microdata DB shows the percentage growth in the last 6 months of Italian companies, spanning various sectors, with a different number of employees, different composition in revenue (residential viz. export) and a different percentage of export to Germany. The attributes of the microdata DB are the direct identifier $\bar{i} = \{Id\}$, where Id is a unique identifier for a company, the quasi-identifiers $\bar{q} = \{Area, Sector, Employees, ResidentialRevenue, ExportRevenue\}$, the non identifying attributes $\bar{a} = \{ExportToDE, Growth6mos\}$, and the weight $W = \{Weight\}$.

Re-identification risk is highest for tuple 15 (0.03) and lowest for tuple 7 (0.003). Extending to general (re-identification-based) statistical disclosure risk, we have that $\rho_{\hat{q}}$ of a given tuple clearly coincides with re-identification risk if \hat{q} includes the Id . It is also the case when \hat{q} includes an n-uple of quasi-identifiers that happens to be unique. For example, tuple 4 is the only one located in the North, dealing in the Textiles sector, with more than 1000 employees; therefore, its re-identification and statistical disclosure risk coincide and amount to 0.016. Notice that its weight (60) witnesses the presence of multiple companies in the identity oracle having the same characteristics as tuple 4 according to the similarity function ϕ , e.g., the same/similar quasi-identifiers.

In another perspective, we are outlining a possible *attack strategy* to attempt re-identification of a given tuple t (Figure 2): 1. filter out a set of tuples C from O that match t on the values of attributes in \bar{q} ; 2. choose the tuple $r \in C$ that best fits t w.r.t. the other attributes; 3. return r with an associated probability/score. To put the attack strategy into action, the entire toolbox from the *record linkage* literature can be adopted [13]. Efficient record linkage techniques typically operate in two steps: *blocking*, when restricting the cohort of candidate matches (step 1 of the attack strategy); *matching*, when evaluating the actual correspondences (step 2). Anonymization techniques aim at making blocking computationally expensive, by suppressing or modifying (as we shall see) selective values, which would make blocking effective restricting the cluster of candidate matches. With large clusters, exhaustive comparison is both computationally expensive, and yields an overly uncertain result, making the attack ineffective.

It is interesting to observe that the sampling weights can be used as a predictor of the effectiveness of a re-identification attack: tuples with higher weights in M will be in clusters with more candidates and thus less likely be identified, though statistically relevant; tuples with lower weights will be in smaller clusters and then will be more easy to re-identify. This gives an optimistic angle on the problem, as anonymization techniques can try to

operate on less representative tuples so as to increase overall confidentiality without hampering the statistical significance.

3 VADALOG REASONING

VADA-SA, the statistical disclosure control framework we introduce in this paper, is based on the VADALOG system, a state-of-the-art *logic-based reasoner* [6] whose core revolves around the VADALOG language, a member of the Datalog[±] family [12, 23]. The disclosure risk measurement techniques as well as the anonymization logic are expressed in VADALOG.

Datalog[±] generalizes Datalog with existential quantification in the rule conclusion, making it suitable for ontological reasoning. A *rule* is a first-order sentence of the form $\forall \bar{x} \forall \bar{y} (\varphi(\bar{x}, \bar{y}) \rightarrow \exists \bar{z} \psi(\bar{x}, \bar{z}))$, where φ (the *body*) and ψ (the *head*) are conjunctions of atoms. For brevity, we omit universal quantifiers and denote conjunction by comma. As usual in this context, the semantics of a set of rules is operationally defined by the well-known CHASE procedure. Intuitively, the CHASE satisfies the rules by generating new head facts for bindings of the body, possibly introducing new variable symbols in the data, in the form of *labelled nulls*, in the presence of existentially quantified head variables [21].

The core of VADALOG is based on *Warded Datalog[±]* [6], a syntactic restriction to Datalog[±] that guarantees decidability and tractability in the presence of recursion and existential quantification. In terms of expressive power, Warded Datalog[±] captures full Datalog and OWL 2 direct semantics entailment regime for OWL 2 QL. The language underpinnings are exploited by the reasoner to allow for efficient execution of reasoning tasks. VADALOG augments Warded Datalog[±] with supplementary features such as aggregation, algebraic operations, and stratified negation. As we shall see, VADALOG is sufficiently expressive to support our anonymization reasoning scenarios and comprises all the needed features such as joint use of full recursion, existential quantification and aggregation, to model propagation of the disclosure risk and anonymize values. These requirements are not met by the standard relational/SQL systems, which in particular offer inefficient or no support for recursion and existentials.

4 THE VADA-SA FRAMEWORK

While statistical disclosure control has traditionally followed a procedural approach, we propose a shift towards a fully declarative one, and look at state-of-the-art reasoners, leveraging our experience on VADALOG and Knowledge Graphs applied to different problems in the financial realm e.g., link prediction [2] as well as schema-independent approaches to model management [3]. The VADA-SA framework, whose architecture is sketched in Figure 3, lies on the following basic pillars:

- A structuring of the statistical disclosure control process in the form of an *anonymization cycle*.
- The construction of a VADALOG-based enterprise *Knowledge Base* (KB) encompassing the patterns and techniques for statistical disclosure risk assessment and anonymization as well as all the surrounding business knowledge to be leveraged.
- The formulation of risk assessment and anonymization phases in the form of *reasoning tasks* upon the KB. In such reasoning tasks, the *extensional component* comprises the microdata DB as well as their basic metadata, such as schema-level information. Much care is devoted to the *intensional component*, encoding reasoning rules for: attribute categorization, risk assessment and anonymization. The intensional component is at high level of abstraction, composed of pluggable VADALOG modules, some of which are provided off-the-shelf while others can

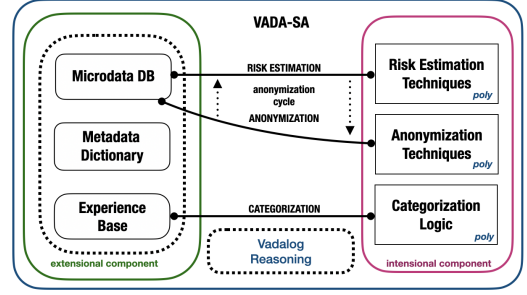


Figure 3: The VADA-SA architecture.

be autonomously developed by business experts. The overall statistical disclosure control process is a reasoning task itself, which relies on the mentioned ones and adaptively chooses the actions to be performed. The *derived extensional component*, i.e., the results of the reasoning process, contains the outcome of risk analysis and the anonymized microdata DBs.

In this section, we first illustrate the anonymization cycle and the *metadata dictionary*, at the basis of our schema-independent approach (Section 4.1), we then focus on the evaluation of statistical disclosure risk (Section 4.2) and anonymization (Section 4.3). Finally, we show extensions and advanced applications involving complex business knowledge (Section 4.4).

4.1 The Anonymization Cycle and the Metadata Dictionary

When a microdata DB needs to be shared, it undergoes the *anonymization cycle* at the core of the VADA-SA architecture, shown in Figure 3. It consists of an iterative application of disclosure risk evaluation and anonymization until the risk is under a given threshold. Each iteration removes a minimum amount of information, and checks whether confidentiality requirements are fulfilled, in a statistics-preserving fashion.

In particular, *risk evaluation* takes as input a microdata DB. Based on its category, each attribute has a different treatment. Direct identifiers must not be disclosed and non-identifying attributes are not needed in the risk evaluation process, thus both are dropped. Quasi-identifiers and the sampling weight are used for disclosure risk estimation. Anonymization is activated until the disclosure risk is acceptable. In so doing, we aim at a trade-off between *statistical preservation* and disclosure risk, as captured by the threshold T , determined on the basis of user experience.

Metadata Dictionary and Attribute Categorization. In order to achieve schema and data independence, in VADA-SA we follow a *meta-level approach* and include a *metadata dictionary* in the KB. Facts of the form `MicroDB(name)`, `Att(microDB, name, description)`, `Category(microDB, att, cat)` are used to reason upon microdata DBs, their attributes and their categories, respectively. Figure 4 shows the portion of VADA-SA dictionary for the “I&G” (Inflation and Growth) microdata DB. Facts for `MicroDB` and `Attribute` are part of the extensional component and change when new microdata DBs are added into VADA-SA. Facts for `Category` are part of the derived extensional component: they are the product of a reasoning process that, for each microdata DB and each attribute, infers the most suitable category. In fact, before entering the anonymization cycle, the attributes of the microdata DB need to be categorized as identifiers, quasi-identifiers or non-identifying attributes, as we have seen in Section 2.1.

Algorithm 1 Attribute categorization

- (1) $\text{Att}(M, A) \rightarrow \exists C \text{ Cat}(M, A, C).$
- (2) $\text{Att}(M, A), \text{ExpBase}(A_1, C), A \sim A_1 \rightarrow \text{Cat}(M, A, C).$
- (3) $\text{Cat}(M, A, C) \rightarrow \text{ExpBase}(A, C).$
- (4) $\text{Cat}(M, A, C_1), \text{Cat}(M, A, C_2) \rightarrow C_1 = C_2.$

Attribute		
Microdata DB	Attribute Name	Description
I&G	Id	Company Identifier
I&G	Area	Geographic Area
I&G	Sector	Product Sector
I&G	Employees	Num. of employees
I&G	Residential Rev.	Rev. from internal market
I&G	Export Rev.	Rev. from external market
I&G	Export to DE	Rev. from DE market
I&G	Growth	Rev. growth last 6 mths
I&G	Weight	Sampling Weight
Category		
Microdata DB	Attribute Name	Category
I&G	Id	Identifier
I&G	Area	Quasi-identifier
I&G	Sector	Quasi-identifier
I&G	Employees	Quasi-identifier
I&G	Residential Rev.	Quasi-identifier
I&G	Export Rev.	Non-identifying
I&G	Export to DE	Quasi-identifier
I&G	Growth	Quasi-identifier
I&G	Weight	Sampling Weight

Figure 4: Metadata Dictionary: Attribute and Category.

Algorithm 1 shows the VADALOG program adopted for this purpose. It features a recursive application of experience: `ExpBase` is the extensional component and stores for an attribute name A , a known category C , according to available experts' knowledge. Assuming one category per attribute (Rule 1), if our attribute is sufficiently similar (according to a pluggable set of similarity functions or denoted by the \sim symbol) to another attribute A_1 of the experience base for which the category is known, we borrow that category (Rule 2), and recursively feed the conclusion back into the experience base (Rule 3), so as to aid other decisions. Rule 4, technically an equality-generating dependency (EGD), guarantees that each attribute is assigned one single category. This VADA-SA module lends itself to human-in-the-loop intervention in two points: when deciding whether to consolidate a decision of Rule 2 with Rule 3, as the user may consider a decision to be use-case specific, and when violations of EGD 4 arise, to allow for manual inspection of doubtful cases.

Anonymization Cycle. The interplay between evaluation of statistical disclosure risk and anonymization is at the core of our framework. Given the input microdata DB, the direct identifiers are removed and all the potentially harmful combinations of quasi-identifiers are evaluated to take countermeasures.

Algorithm 2 Anonymization cycle

- (1) $\text{Val}(M, I, A, V), \text{Cat}(M, A, C), C \in \{\text{Quasi-identifier}, \text{Weight}\}, \text{VSet} = \text{munion}((A, V)) \rightarrow \text{Tuple}(M, I, \text{VSet}).$
- (2) $\text{Tuple}(M, I, \text{VSet}), \# \text{risk}(I, R), R > T \rightarrow \# \text{anonymize}(I).$
- (3) $\text{Tuple}(M, I, \text{VSet}), \# \text{risk}(I, R), R \leq T \rightarrow \text{Tuple}_A(M, I, \text{VSet}).$

The set of VADALOG rules of Algorithm 2 compactly represents this logic. Rule (1) creates `Tuple` facts for each tuple of the microdata DB M , identified by an artificial identifier I , and collects all the name-value pairs for quasi-identifiers and sample weights into the `VSet` variable. `Val` facts are part of the extensional component and store the value V for an attribute A of the microdata DB M . The identifiers of M are implicitly dropped. Observe that *munion* performs such aggregation, for each microdata DB M and `VSet` is a set-type variable. Whenever a specific tuple I violates a $[0, 1]$ risk threshold T , a fact `anonymize` is produced for I , in Rule 2. Both `risk` and `anonymize` are atoms defined in external libraries, in VADALOG (denoted by the “#” prefix). In particular, `risk` returns the risk R associated to a given tuple I ; it is a compact form for the join “`#riskInput(I), #riskOutput(I, R)`”, where `riskInput` is a fact triggering a VADALOG program producing facts of `riskOutput` for I . More simply, `anonymize` produces new facts for `Tuple`. This mechanism embeds a recursion on Rule 2, to anonymize tuples that still do not pass the risk evaluation. Only those facts for `Tuple` that pass the risk validation of Rule 3 are copied to `TupleA`, which can be considered anonymized.

The anonymization cycle in Algorithm 2 makes the approach *fully explainable* in the sense that each anonymization decision taken by Rule 2 is motivated by the specific binding of its body. It is also *preemptive* and *active*, in the sense that for each threshold violation, greedily applies a single anonymization step, at the same time minimizing the amount of suppressed statistical information. It is *schema independent*, as only atoms of the meta-data dictionary are used and there is no specific reference to either schema or instance objects of the single microdata DBs. We shall see how specific values are bound to the attributes as a responsibility of risk and anonymize implementations. Finally, the algorithm offers multiple degrees of freedom: different risk and anonymization techniques can be used, and our `risk` and `anonymize` and polymorphic, in this sense; specific optimizations and execution heuristics can be adopted to choose which tuples to anonymize first (by controlling the activation order of Rule 2 against its possible bindings), and to choose which quasi-identifiers to anonymize first.

4.2 Statistical Disclosure Risk Estimation

Our risk atom in Algorithm 2 is polymorphic. VADA-SA features a plug-in mechanism to opt for specific implementations at runtime. While the high-level characteristics of the VADALOG language allow to delegate users to specify their own risk logic, with extensive use of business knowledge, a number of techniques are provided off-the-shelf. In this section, we introduce the main risk disclosure evaluation techniques offered by VADA-SA.

Re-identification-based. We start in Algorithm 3 with re-identification-based risk evaluation, that we have defined in Section 2.2.

Algorithm 3 Re-identification-based risk evaluation

- (1) $\text{Tuple}(M, I, \text{VSet}), \text{riskInput}(I), \text{Cat}(M, W, \text{Weight}), R = 1/\text{msum}(\text{VSet}[W], \langle I \rangle) \rightarrow \text{Tuple}_A(R, \# \text{VSet}[\text{AnonSet}]).$
- (2) $\text{Tuple}(M, I, \text{VSet}), \text{Tuple}_A(R, \text{VSet}^*) \rightarrow \text{riskOutput}(I, R).$

Whenever a tuple I needs to be evaluated (`riskInput` atom), in Rule 1, the name W of the weight atom is retrieved from the metadata dictionary and used to extract the weight value from `VSet`, with the *access* operator denoted by $[X]$, where X can be either a single attribute name or a collection thereof. Weights are summed (*msum*) and the risk score $1/R$ is computed. `TupleA` has

variable arity, and its terms are the computed risk R and the set of quasi-identifiers to group by when forming the summation. The expression $*VSet[AnonSet]$ has the following meaning: the prefix operator “*” is *collection unpacking* and turns each element of the argument collection into a term of $TupleA$, essential for grouping along quasi-identifiers. Note that $VSet$ is filtered by a set $AnonSet$ of attribute names—the order is irrelevant in this context—that selects those that are considered of interest by business experts w.r.t. risk evaluation. Rule 2 finds those tuples I for which the risk has been computed and returns it. It uses the *packing operator*, denoted by the suffix operator “*”, which packs a sequence of terms of $TupleA$ into the set variable $VSet$. In this way, by joining on $VSet$, we can identify all the tuples to which the risk computation applies. We have already seen examples of re-identification-based risk estimation in Section 2.2.

k-anonymity is a commonly used threshold approximation of re-identification risk estimation [34]. For a given set of quasi-identifiers, whenever the number of occurrences is less than a fixed threshold k , it is considered dangerous; it is safe otherwise. For instance, in the microdata DB of Figure 1, considering the quasi-identifiers *Area* and *Sector*, we notice, e.g., that there is only one occurrence for “North” and “Public Service” (tuple 1). We say the set of pairs $\{\langle Area, North \rangle, \langle Sector, Public Service \rangle\}$ is a *sample unique* for tuple 1. The VADALOG reasoning rules for k-anonymity are reported in Algorithm 4.

Algorithm 4 k-anonymity

- (1) $Tuple(M, I, VSet), riskInput(I),$
 $R = mcount(\langle I \rangle) \rightarrow TupleA(R, *VSet[AnonSet]).$
 - (2) $Tuple(M, I, VSet), TupleA(R_1, VSet*)$
 $R = case R_1 < k then 1 else 0 \rightarrow riskOutput(I, R).$
-

Individual Risk. In the re-identification model the simplifying assumption is made that the sampling weight W_t corresponds to the frequency (number of occurrences) F_k of a given combination k of quasi-identifiers in the total population from which the microdata DB has been sampled; therefore we can compute the combination risk as $1/F_k$. Yet, frequencies F_k are unknown and in general different from W_t . A further inferential step is then required. The typical approach [7, 22, 38] is accounting for F_k in a Bayesian fashion, by considering the distribution of the population frequencies given the sample frequencies $F_k|f_q$ and obtaining $1/F_k$ as the posterior mean. In our setting, the sample frequency is the sample count in the microdata DB. Different assumptions can be made on the posterior distribution of $F_k|f_q$, with different techniques to accordingly estimate ρ_q . The one we adopt here is considering such distribution a negative binomial and thus we pose $\lambda = \sum W_t/f_q$ to estimate risk in Equation 1 of Section 2.2. Indeed, other distributions can be adopted. The individual risk estimation is formalized in Algorithm 5.

Algorithm 5 Individual risk

- (1) $Tuple(M, I, VSet), riskInput(I), Cat(M, W, Weight),$
 $F = mcount(\langle I \rangle), R = msum(VSet[W], \langle I \rangle) \rightarrow$
 $TupleA(F/R, *VSet[AnonSet]).$
 - (2) $Tuple(M, I, VSet), TupleA(R, VSet*) \rightarrow riskOutput(I, R).$
-

While scanning through the tuples having a given combination $VSet$, used as a group-by key thanks to the unpacking operator, Rule 1 counts the occurrences of each combination (frequency) and sums the contributor weights. Facts for $TupleA$ are produced

only once all the contributors are available. The risk is then estimated for each combination and finally returned by Rule 2.

SUDA. With k-anonymity, we have introduced the concept of sample unique, i.e., a set of quasi-identifiers—name-value pair—that identify a tuple of a given microdata DB, i.e., they are unique. A sample unique is not the same as a database key, because it expresses a property that holds at tuple level and not at schema level. Alongside the schema-level distinction between superkey and key (a minimal superkey) in relational theory [1], here, at data level, we introduce the *minimal sample unique* (MSU) μ_t for a given tuple, that is a sample unique for which there exists no other sample unique μ'_t for the same tuple, s.t. $\mu'_t \subset \mu_t$. The *Special Unique Detection Algorithm* (SUDA) is a heuristic technique that estimates the statistical disclosure risk of a given tuple based on the size and the number of its MSUs.

Consider for example the set $\mu_{20}^1 = \{\langle Area, Center \rangle, \langle Sector, Financial \rangle, \langle Employees, 1000+ \rangle, \langle Res. Rev., 30-60 \rangle\}$ for the microdata DB in Figure 1 for tuple 20. It is sample unique though not MSU, since the set $\mu_{20}^2 = \{\langle Sector, Financial \rangle\}$ is sample unique and s.t. $\mu_{20}^2 \subset \mu_{20}^1$. Moreover, μ_{20}^2 is MSU. Similarly, $\mu_{20}^3 = \{\langle Employees, 1000+ \rangle, \langle Res. Rev., 30-60 \rangle\}$ is another MSU. In total, tuple 20 has 2 MSUs.

Algorithm 6 encodes the VADA-SA version of SUDA.

Algorithm 6 SUDA

- (1) $Tuple(M, I, VSet), riskInput(I) \rightarrow TupleI(M, I, VSet).$
 - (2) $TupleI(M, I, VSet), Cat(M, A, Quasi-identifier),$
 $A \in VSet \rightarrow \exists Z Comb(Z, I), In(A, Z).$
 - (3) $Comb(Z_1, I), TupleI(M, I, VSet), Cat(M, A, Quasi-identifier),$
 $A \in VSet, not In(A, Z_1) \rightarrow$
 $\exists Z Comb(Z, I), InComb(Z, Z_1), In(A, Z_1).$
 - (4) $InComb(X, Y), In(A, X) \rightarrow In(A, Y).$
 - (5) $Comb(Z, I), In(A, Z), TupleI(M, I, VSet),$
 $ASet = munion(A) \rightarrow TupleC(I, *VSet[ASet]).$
 - (6) $TupleC(I, VSet*), mcount(\langle I \rangle) = 1 \rightarrow$
 $\exists S Su(S, VSet), HasSu(I, S).$
 - (7) $Su(S, VSet), HasSu(I, S), not HasSu(I, S_1),$
 $Su(S_1, VSet'), VSet' \subset VSet \rightarrow MSU(I, S).$
 - (8) $TupleI(M, I, VSet), MSU(I, S), Su(S, VSet),$
 $R = case size(VSet) < k then 1 else 0 \rightarrow riskOutput(I, R).$
-

After restricting the focus on input tuples (Rule 1), for each tuple we generate all the combinations of quasi-identifiers, first by introducing a combination Z for each of them (Rule 2), and then by constructing all the possible extensions that can be obtained by adding other quasi-identifiers (Rules 3 and 4). Then, for each combination of quasi-identifiers, generated by unpacking (Rule 5), we generate sample unique facts for Su , denoting those combinations that exactly identify one single tuple. The predicate $HasSu$ is needed since every tuple I can have multiple sample unique sets, while the *mcount* aggregation needs to group by $VSet$. Rule 7 creates facts for MSU, filtering only those sample unique sets that are minimal. Finally, Rule 8 implements the logic to handle minimal sample unique sets. In this case, we evaluate the size of every MSU and if it is above a given threshold k , we consider the input tuple dangerous and thus return 1. The assumption here is that we cannot accept that the number of quasi-identifiers that can disclose the identity is too small. Clearly, more sophisticated checks could be implemented, possibly also including an overall evaluation of all the MSUs for a given tuple, for example by comparing the average size of MSUs against a threshold.

	I	Q	Q	Q	Q	F
	Id	Area	Sector	Employees	Residential Revenue	
1	099876	Roma	Textiles	1000+	0-30	1
2	765389	Roma	Commerce	1000+	0-30	2
(a) 3	231654	Roma	Commerce	1000+	0-30	2
4	097302	Roma	Financial	1000+	0-30	2
5	120967	Roma	Financial	1000+	0-30	2
6	232498	Milano	Construction	0-200	60-90	1
7	340901	Torino	Construction	0-200	60-90	1

	I	Q	Q	Q	Q	F
	Id	Area	Sector	Employees	Residential Revenue	
1	099876	Center	\perp_1	1000+	0-30	5
2	765389	Center	Commerce	1000+	0-30	3
(b) 3	231654	Center	Commerce	1000+	0-30	3
4	097302	Center	Financial	1000+	0-30	3
5	120967	Center	Financial	1000+	0-30	3
6	232498	North	Construction	0-200	60-90	2
7	340901	North	Construction	0-200	60-90	2

Figure 5: Local suppression and global recoding.

4.3 Smart Anonymization

In Algorithm 2 we have introduced the anonymization cycle, where Rules 2 and 3 show the interaction between risk estimation, with the main techniques introduced in Section 4.2, and anonymization, the object of this section. Tuples whose risk is considered over a given threshold T , produce facts for anonymize, which polymorphically triggers dedicated VADALOG programs: for each tuple I having statistical disclosure risk $R > T$, a new fact for Tuple is produced, with the same identifier I and statistical disclosure risk $R' < R$. The process continues recursively, until there is no tuple violating the threshold.

All the statistical disclosure evaluation techniques of Section 4.2 compute the risk associated to each tuple with monotonic aggregations, which play an important role here. In particular, all of them (e.g., *msum* in Algorithm 3, *mcount* in Algorithm 4, etc.) take as input the *aggregation contributor*, denoted by $\langle I \rangle$. According to the monotonic aggregation semantics [6], whenever two or more tuple tuples having the same value for the contributor I are aggregated (e.g., summed, counted, etc.) within the same group (defined by the bindings of head variables), only the tuple providing the least risk contribution is considered, while the others are neglected. This implies that, whenever a tuple I is replaced by a “more anonymous version”, for example by suppressing a quasi-identifier, as we shall see, as the two are seen as the same contributor (they have the same value for I), only the anonymized one will be accounted for in the aggregation, so that more anonymized tuples incrementally replace the others and reduce risk, until convergence is achieved. We anonymize tuples with two main techniques: *introducing labelled nulls* to replace selective values, *applying a global recoding*.

Local Suppression with Labelled Nulls. Labelled nulls are a powerful tool from logic-based reasoning, which we effectively apply in the anonymization context. Consider the microdata DB in Figure 5a, where all the attributes are assumed to be quasi-identifiers, the sampling weight is omitted for simplicity, and the frequency of the n -uple of quasi-identifiers is showed on the right. For tuple 1 the set $\{\langle \text{Area}, \text{Roma} \rangle, \langle \text{Sector}, \text{Textiles} \rangle, \langle \text{Employees}, 1000+ \rangle, \langle \text{Resid. Rev.}, 0-30 \rangle\}$ is sample unique. What if we replace the value “Textiles” for *Sector* with a labelled null \perp_1 ? As we are not aware of the underlying value of \perp_1 , the combination of quasi-identifiers at hand may match with any among tuples 2-5, thus leading to a total frequency of 5. Likewise, tuples 2-5 see their frequency increased to 3. In total, by adding a single labelled null and hence introducing some degree of uncertainty, we have highly decreased the statistical disclosure risk of the microdata DB, as it can be seen in Figure 5b. In fact, a tuple

containing one or more nulls may match with different tuples of the microdata DB, or even with none of them, depending on the specific assignment for those nulls.

Going back to what introduced in Section 2.2, in our framework in order to estimate the statistical disclosure risk, we need to compute $\lambda(\sigma_{\hat{q}=\hat{q}} M)$ over a selection of the microdata DB M , based on an n -uple of values \hat{q} for quasi-identifiers. The risk estimation techniques of Section 4.2 apply λ to the entire microdata DB and the selection is implicit in the grouping performed by the aggregations, in the sense that, for each tuple, the aggregation forms the group by selecting only those tuples having the same values for quasi-identifiers (or subset of interest, thereof). So, $M(\tilde{i}, \tilde{q}, \dots)$ is included in the selection induced by tuple $M(\tilde{i}, \tilde{q}, \dots)$ iff $(q'_1, \dots, q'_n) = (q_1, \dots, q_n)$, and, by construction, the groups form a partition of the microdata DB. If we allow q_i to be a labelled null, a new semantics must be adopted to define whether $q_i = q'_i$ and thus form the aggregation groups.

The introduction of nulls raises non-trivial semantic issues when aggregations are involved, and theoretical work is still needed to achieve sound characterizations [25]. In VADA-SA, for the construction of aggregation groups, we adopt a *null-tolerant semantics* inspired by the so-called *maybe-match* approach [14], and assume that $q_i =_\perp q'_i$ holds if: (i) q_i and q'_i have the same constant value, or (ii) either q_i or q'_i is a labelled null. Consequently, $(q'_1, \dots, q'_n) =_\perp (q_1, \dots, q_n)$ holds iff $q_i =_\perp q'_i$ holds for every $1 \leq i \leq n$. The $=_\perp$ relation is therefore used, instead of standard equality, to form groups. A tuple containing null quasi-identifiers, as a result of anonymization steps, is assigned to multiple aggregation groups (which do not partition the microdata DB anymore), increasing their cardinality and so anonymity.

We now have all the ingredients ready to encode local suppression, an anonymization method where quasi-identifiers are replaced by labelled nulls to reduce the statistical disclosure risk. The technique is expressed by Algorithm 7.

Algorithm 7 Local suppression

(1) Tuple($M, I, VSet$), anonymize(I), Cat(M, A , Quasi-identifier), $VSet[A]$ is not null $\rightarrow \exists Z$ Tuple($M, I, (A, Z) \cup (VSet \setminus (A, _))$).

For a tuple I that needs to be anonymized, as witnessed by the predicate anonymized, for a not null quasi-identifier A , we generate a new tuple, where it is replaced by a labelled null Z .

Global Recoding. While local suppression introduces nulls, another technique to control statistical disclosure risk consists in decreasing the granularity of the values on the basis of domain knowledge. Consider again Figure 5a. Tuples 6 and 7 have the following sample unique sets, respectively: $\{\langle \text{Area}, \text{Milano} \rangle, \langle \text{Sector}, \text{Construction} \rangle\}, \{\langle \text{Area}, \text{Torino} \rangle, \langle \text{Sector}, \text{Construction} \rangle\}$ and therefore have high disclosure risk. Besides the basic meta-data dictionary we have seen in Section 4.1, the VADA-SA KB contains knowledge about the attribute domains as well as the mutual relationship between their values. For instance, for the attribute *Area*, the KB comprises the following information:

```
Att(I&G, Area). TypeOf(Area, City). SubTypeOf(City, Region).
InstOf(Milano, City). InstOf(Torino, City).
InstOf(North, Region). IsA(Milano, North). IsA(Torino, North).
```

The *Area* attribute is known to be of Type “City”, which in turn is a SubtypeOf “Region”. Moreover we know that Milano and Torino are instances of cities and North is an instance of region. Finally, we now that both Milano and Torino are in the North. Similar knowledge is present for the entire geography.

Algorithm 8 Global recoding

(1) $\text{Tuple}(M, I, \text{VSet})$, $\text{anonymize}(I)$, $\text{Cat}(M, A, \text{Quasi-identifier})$,
 $\text{TypeOf}(A, X)$, $\text{subTypeOf}(X, Y)$, $\text{isA}(\text{VSet}[A], Z)$,
 $\text{TypeOf}(Z, Y) \rightarrow \text{Tuple}(M, I, (A, Z) \cup (\text{VSet} \setminus (A, _)))$.

The logic for global recoding is in Algorithm 8. For a tuple that needs to be anonymized, we consider a quasi-identifier A . Based on its type, we climb the hierarchy up to its direct super-type Y . Then for the value $\text{VSet}[A]$ of A , we use the corresponding value Z of Y to replace $\text{VSet}[A]$. This form of suppression can be effectively applied to the entire microdata DB (and in this sense it is “global”) and is inherently recursive as multiple hierarchical roll-ups may be needed to guarantee anonymity.

4.4 Enhancing Anonymization

We conclude the section by discussing two advanced topics: embedding of complex business knowledge, where we showcase the use of domain experience for *context aware* anonymization, and implementation of runtime heuristics, to maximize the statistical effectiveness of our approach.

Embedding complex business knowledge. The overall anonymization process can largely benefit from the surrounding business knowledge, an aspect often neglected by dedicated tools. Thanks to reasoning, we can inject business representations into different phases of Algorithm 2: in risk estimation modules, to craft ad-hoc methods; into anonymization techniques, e.g., to opt for specific values for global recoding, and so on. The setting we show here, motivated by our experience in the Bank of Italy with financial networks, consists of taking into account the relationships that exist between the respondents, say X and Y . It is in fact common that the statistical disclosure risk propagates along linked entities, e.g., companies or people, so that being able to re-identify one, makes it easier to re-identify others. In essence, all the linked entities of a given cluster, have the same disclosure risk, obtained as the probability that at least one entity of the cluster is re-identified: $1 - \prod_c (1 - \rho^c)$, where ρ^c is the risk of an entity, calculated with one of the techniques in Section 4.2. Here, along the lines of what usually done to estimate the risk of *households* and *hierarchical structures* [26], re-identification risk is interpreted as the re-identification probability.

Now, types of links that can be considered are arbitrarily complex: finding members of the same family, companies of the same company group are examples. The latter, e.g., could be encoded by the following VADALOG rules: (1) $\text{Own}(X, Y, W)$, $W > 0.5 \rightarrow \text{rel}(X, Y)$. (2) $\text{rel}(X, Z)$, $\text{Own}(Z, Y, W)$, $\text{msum}(W, \langle Z \rangle) > 0.5 \rightarrow \text{rel}(X, Y)$. Clusters of companies ($\text{rel}(X, Y)$ holds where X and Y are in the same cluster) are defined by *company control relationships*: if X owns more than 50% of the shares of Y (Rule 1) or controls a set of companies Z that jointly own more than 50% of Y , than X controls Y and thus X and Y are in the same cluster.

Algorithm 9 shows an enhanced version of Algorithm 2, where the risk for a tuple I_2 is estimated as explained. Specifically, Rule 2 uses $\#rel$ (and we assume here $\text{rel}(X, X)$ holds) to compute the risk for I_1 as the combined risk of the entities in the same cluster. The aggregation *mprod* is the monotonic product, which considers, for each contributor I_2 , the maximum contribution it provides, so as to account for the new less risky anonymized tuples produced by Rule 3, and eventually triggering Rule 4.

Runtime heuristics. We have seen how VADA-SA operates incrementally and applies anonymization steps, only when tuples exhibit an overly high statistical disclosure risk. However, there

Algorithm 9 Enhanced anonymization cycle

(1) $\text{Val}(M, I, A, V)$, $\text{Cat}(M, A, C)$, $C \in \{\text{Quasi-identifier}, \text{Weight}\}$,
 $\text{VSet} = \text{munion}((A, V)) \rightarrow \text{Tuple}(M, I, \text{VSet})$.
(2) $\text{Cat}(M, A, C)$, $C = \text{Identifier}$, $\text{Tuple}(M, I_1, \text{VSet}_1)$,
 $\text{Tuple}(M, I_2, \text{VSet}_2)$, $\#rel(\text{VSet}_1[A], \text{VSet}_2[A])$, $\#risk(I_1, R)$,
 $R_{\text{clust}} = 1 - \#mprod(1 - R, \langle I_2 \rangle) \rightarrow \text{Risk}(I_1, R_{\text{clust}})$.
(3) $\text{Tuple}(M, I, \text{VSet})$, $\text{Risk}(I, R)$, $R > T \rightarrow \#anonymize(I)$.
(4) $\text{Tuple}(M, I, \text{VSet})$, $\text{Risk}(I, R)$, $R \leq T \rightarrow \text{Tuple}_A(M, I, \text{VSet})$.

are still various open questions to be addressed, which correspond to specific degrees of freedom in anonymizing microdata. If there are two or more tuples that violate the risk threshold, which ones should be anonymized first? Moreover, if there are two or more quasi-identifiers of the same tuple, which one should be suppressed or recoded first?

As for the first question, in VADA-SA, we adopt a *greedy* approach and choose to anonymize first the tuples that carry less statistical significance (namely, *data utility*), which can be estimated on the basis of the sampling weight. We exploit the so-called *routing strategies* [5] of the underlying VADALOG system to decide which bindings of the rule body to privilege when multiple possibilities arise. The approach here is quite intuitive: a “*less significant first*” strategy sorts the bindings of Rule 2 by risk and guides the anonymization accordingly.

The second question, namely the prioritization of quasi-identifiers, requires more care. We have seen in Algorithms 7 and 8 that either an existential or a higher-level domain value is used to replace quasi-identifiers and the specific attribute to consider is chosen as a consequence of the binding of $\text{Cat}(M, A, \text{Quasi-identifier})$. Also in this case, we can prioritize bindings by adopting a VADALOG routing strategy and the greedy approach. In particular, a “*most risky first*” strategy would first bind the rules against the attributes that affect more the tuple-level disclosure risk. So, in this case, the strategy itself would rely on a VADALOG program computing the risk, in order to take informed decisions. For instance, consider the problem of anonymizing tuple 1 of Figure 5a. Applying local suppression on *Sector* removes any sample unique of the tuple, which then occurs with frequency 5; instead, applying local suppression on *Area*, e.g., would leave the value “Textiles” for *Sector*, and would then require further local suppressions until such attribute is removed, with a consequential loss of data utility. In other terms, a greedy approach to local suppression or global recoding sustains the preservation of data utility.

5 EXPERIMENTS

VADA-SA has been fully implemented and engineered in the VADALOG System. Towards a production application of the framework for the Research Data Center of the Bank of Italy, the system has been extensively experimented on real-world datasets from the Bank of Italy and synthetic ones to assess its *anonymization capability* (Section 5.1) and *scalability* (Section 5.2). The schema independent approach makes the framework general purpose and suitable for treating datasets in any domain.

Datasets. The microdata DBs used in the experimental analysis are reported in Figure 6. The *real-world* and *realistic datasets* derive from the *Inflation and Growth Survey* of the Bank of Italy, whose schema is shown in Figure 1. The *synthetic datasets* have been generated by fitting the real-world distribution (denoted by “W” in the figure) or by inducing specific unbalanced or very unbalanced distributions (denoted by “U” and “V”). Unbalanced

Dataset	No. Att.	No. Tuples	Dist.	Data
R6A4U	4	6k	U	Synth
R12A4U	4	12k	U	Synth
R25A4W	4	25k	W	Real-world
R25A4U	4	25k	U	Realistic
R25A4V	4	25k	V	Realistic
R50A4W	4	50k	W	Synth
R50A4U	4	50k	U	Synth
R50A5W	5	50k	W	Synth
R50A6W	6	50k	W	Synth
R50A8W	8	50k	W	Synth
R50A9W	9	50k	W	Synth
R100A4U	4	100k	U	Synth

Figure 6: Datasets used in the experimental settings.

distributions comprise many tuples with very selective combinations of quasi-identifiers, which exhibit high disclosure risk.

Hardware. We employed a memory-optimized virtual machine with 16 cores and 128 GB RAM on an Intel Xeon architecture.

5.1 Testing Anonymization Capability

We analyzed the capability of the system to detect tuples that need to be anonymized (i.e., the “risky” tuples).

Reduction of risk vs. loss of information. We applied the VADA-SA anonymization cycle to the real-world dataset *R25A4W* with the *k*-anonymity risk evaluation technique (Section 4.2) and choosing a risk threshold $T = 0.5$. We employed *local suppression* anonymization (Section 4.3), with a *less significant first* runtime heuristic (Section 4.4). We varied the anonymity threshold k from 2 to 5. We adopted two metrics to evaluate the capability of the system to detect risky tuples: we counted the number of nulls injected by the local suppression as a result of risk evaluation, so analyzing how many values the system was able to erase (Figure 7a); we estimated the *loss of information* by weighing the number of erased values (i.e., the injected nulls) by the maximum total number of values, those of quasi-identifiers of the risky tuples w.r.t. T , that can be theoretically removed (Figure 7b) to satisfy the *k*-anonymity requirement. For both the measurements, we also evaluated the robustness of the approach, by applying the anonymization cycle to artificial but realistic datasets (*R25A4U* and *R25A4V*) having the same distribution of quasi-identifiers of *R25A4W*, but with an increased number of risky tuples.

The results in Figure 7a confirm what expected: when the *k*-anonymity threshold is increased, the anonymization cycle becomes less tolerant, and more redundancy is required to guarantee anonymity, therefore, in absolute terms, more and more labelled nulls need to be added to suppress specific values. We also observe that the number of needed nulls linearly grows with the tolerance threshold, as a consequence of an overall uniform distribution of values in the adopted combination of 4 quasi-identifiers. While an average real-world dataset requires less than 50 labelled nulls for 25k tuples with a 5-tuples tolerance threshold, more unbalanced versions require more, while confirming the trend. Figure 7b witnesses very good behaviour of VADA-SA in terms of statistical preservation. For the real-world and the mildly unbalanced dataset, the loss of information is constantly below 20%, in particular between 12% (lower bound of *R25A4W*) and 17% (upper bound of *R25A4U*). For these two datasets, the constant trends show that when the number of risky tuples increases, the greedy approach succeeds in removing the values with a wider risk reduction effect. The loss of information for the very unbalanced dataset *R25A4V* is clearly higher, 37%, but it interestingly drops to 13% with less tolerant runs, because

the high number of tuples that are considered risky on different combinations of values of quasi-identifiers, collapse in the *k*-anonymity comparison, when labelled nulls are introduced: so, while the number of nulls is high in absolute terms, the loss of information decreases. This result turns out to be an extremely positive guarantee of the anonymization capability of VADA-SA.

Maybe-matching labelled nulls. In this experiment, we want to assess the effectiveness of the maybe-match semantics, which we use to compare labelled nulls with one another (and has been described in Section 4.3), as opposed to the standard semantics of labelled nulls, such as that adopted in CHASE-based procedures (e.g., Skolem CHASE [11]). According to the standard semantics, for a quasi-identifier q_i , we have that $q_i =_{\perp} q'_i$ holds if: (i) q_i and q'_i have the same constant value, or (ii) both q_i and q'_i are labelled with the same null symbol. We plugged this semantics into VADA-SA, used the same real-world and realistic datasets of Figure 7b and report the number of injected nulls by *k*-anonymity threshold in Figure 7c. Also here, the risk threshold $T = 0.5$ has been used. The figure highlights the proliferation of symbols (the red lines) that takes place with the standard semantics, which is in fact unusable in this setting. By contrast, the probabilistic interpretation of nulls we foster, minimizes the number of labelled nulls (the light-blue lines, whose zoomed version is in Figure 7a).

Using business knowledge. We show the results of anonymization in a real-world setting where anonymization cycle is complemented with a set of VADALOG rules that produce derived extensional knowledge about control relationships between companies. The rules and the setting have been presented in detail in Section 4.4. For the test, we adopt the real-world dataset *R25A4W* and its tweaked unbalanced versions, *R25A4U* and *R25A4V*. We anonymize each of the datasets by estimating the risk with *k*-anonymity with $k = 2$ and $T = 0.5$. We measure the number of nulls injected by local suppression in 5 settings, with increasing number of inferred control relationships, from 0 to 400.

The results are shown in Figure 7d. With all the datasets, the number of injected nulls grows with the number of relationships between entities, which induce bigger and more risky clusters. The three distributions of the quasi-identifier values differently interact with the derived relationships: the more unbalanced the dataset is, the more tuples will be affected by the propagation of risk of the outliers, resulting into a globally risky dataset, to be severely anonymized. In real-world tests, relationships disclose many cases that deserve anonymization (from 9 in the case of 100 relationships to 38 for 400), while the propagation effect is maximized in the *R25A4V* dataset with an upper bound of 323 injected nulls for 300 relationships.

5.2 Testing Scalability

Given the characteristics of the data at hand to be anonymized, we need to make sure that our approach scales well. Although the anonymization cycle, risk estimation and anonymization of VADA-SA are expressed in VADALOG, where reasoning is PTIME in data complexity [6], here we want to investigate on the specific runtime of the system in different settings.

By dataset size. We tested the scalability of VADA-SA by increasing volumes, with 4 synthetic datasets (from *R6A4U* to *R100A4U*), unbalanced and having a high number of risky tuples. We measured the elapsed time for the entire anonymization cycle and also pointed out the sole risk estimation component, with 3 different risk estimation techniques (*individual risk*, *k*-anonymity, *SUDA*). We used $k = 2$ for *k*-anonymity, 3 as the MSU threshold




Figure 7: (a) Number of nulls injected by k -anonymity threshold. (b) Information loss by k -anonymity threshold. (c) Number of nulls injected with maybe-matching viz. standard labelled null semantics. (d) Number of nulls injected by increasing number of relationships in settings with explicit modeling of business knowledge. (e) Execution time by dataset size and risk estimation technique. (f) Execution time by number of quasi-identifiers and risk estimation technique.

for SUDA (see Section 4.3) and $T = 0.5$. We executed each run 5 times and averaged the measurements, for a total of 60 runs.

The results are shown in Figure 7e. All the three groups of trends confirm that the risk estimation component (dotted lines) dominate the elapsed time. This is reasonably expected, as the convergence of our anonymization cycle depends on a positive evaluation of risk estimation, which is then the bottleneck. The linear trend confirms the applicability of the approach. In particular, k -anonymity exhibits a very good behaviour, with elapsed time between 6 and 192 seconds for 100k tuples. The limited cost of estimation can be ascribed to the adoption of monotonic aggregations, which adopt incremental updates and need not be recomputed from time to time. Whilst in Algorithm 5 we have made a simple assumption to estimate the risk from the posterior distribution $F_k | f_{\hat{q}}$ (which would have led to elapsed times similar to those for k -anonymity), for this experiment we plugged into VADA-SA an off-the-shelf statistical library and sampled from the actual negative binomial distribution. The costly trend is motivated by the interaction overhead between the native VADALOG component and the library. The trend for SUDA is less than linear, with, e.g., 727 seconds for 50k and 1344 seconds for 100k tuples since the potential blowup on the number of examined combinations of quasi-identifiers is controlled by the VADALOG optimizations.

By number of quasi-identifiers. To investigate more the dependence of performance on the number of quasi-identifiers, we stressed VADA-SA by anonymizing 6 datasets R50A4W-R50A9W, so with increasing number of attributes and fixed number of tuples, 50k, and real-world-like distribution. We used the same thresholds for k -anonymity, SUDA, and T . We measured elapsed

time and the risk estimation component. We executed each run 5 times averaging the results, for a total of 90 runs.

Figure 7f reports the results. As expected, individual risk and k -anonymity are only marginally affected by the increased number of quasi-identifiers, as they do not consider all the combinations with at most k attributes, but only those with exactly k . Instead, we may expect a much worse trend for SUDA, where for each tuple, all the combinations of at most k attributes are inspected to detect potential MSU. Remarkably, no combinatorial blowup appears in the figure, witnessing a very effective behaviour of the VADALOG execution optimization: while the activation of Rules 2-5 of Algorithm 6 could in theory cause a blowup w.r.t. k , it does not happen in practice because the greedy activation of Rule 7 performed by VADALOG to detect the MSUs preempts the generation of redundant combinations of quasi-identifiers.

6 RELATED WORK

Statistical disclosure control is a broad topic to which many have contributed, especially from the Statistics community, whose work can be considered related to ours.

The concept of *Sample Uniqueness* (SU) to measuring the risk of data disclosure was introduced by Skinner [35], while *k-anonymity*, was first presented by Sweeney [37], along with the first methods of anonymization by generalization (our *global recoding*) and *local suppression*. The measure of *individual risk* in our contribution is inspired by the work of Benedetti and Franci [7] who proposed to compute the risk of data disclosure with the sampling weights of data records.

The topic of data anonymization is related to the area of *differential privacy* [17], where an interesting concept may be adopted in our approach so as to develop a new family of risk measures,

based on the idea that an individual's privacy may be violated even knowing the absence of the individual from the microdata. Investigating such direction will be matter of future work.

While the foundations of our work are set in the theory of statistical disclosure control, our contribution is concerned with providing an industrial production ready solution for the Bank of Italy, conveying a set of properties that derive from a fully declarative reasoning approach. In this sense, we combine our experience in logic-based reasoning [6] and schema-independent solutions to model management problems [3]. None of the existing dedicated software solutions for statistical disclosure control offers the mentioned set of properties. The software pack ARGUS [27] aims at local suppression and coding, as does the Datafly system [36]. Manning et al. introduce SUDA2 (Special Unique Detection Algorithm) [29], whose objective is to detect the risk in certain unique combinations of variables. Recently, the *R* package *sdcMicro* has implemented many of the risk measures and anonymization approaches of our interest [9]. Likewise, ARX is a solution for data anonymization that has been proposed as a practical approach to Statistical disclosure control [33]. A comprehensive survey of the statistical approaches has been provided by Matthews and Harel [30]. Recent work on the risk of information disclosure in linked data, and, more in general, ontology-based data, has formalized the problem and defined its logical foundations [8], with an interest in the concept of linkage safety in RDF graphs [24]; a declarative framework for linked data anonymization has also been proposed [15]. The problem of preserving privacy in data exchange has been analyzed also in the context of information integration systems [31] where a practical solution is represented by *MapRepair* [10] and in the cryptography community, with homomorphic encryption [28].

In the AI literature, statistical disclosure control has been mostly considered within machine learning [16] and deep learning approaches [4]. Yet, they have a different focus and aim at generating anonymized clones of existing datasets while respecting the original statistical properties. An interesting deductive proposal by Øhrn and Ohno-Machado uses Boolean reasoning for data anonymization in databases [32], which however remains purely theoretical and just considers the combinatorial aspect.

7 CONCLUSION

In this paper, we presented VADA-SA, a declarative statistical disclosure control framework. We demonstrated the anonymization workflow, metadata dictionary, and statistical disclosure risk estimation. Utilizing these components, we introduced the anonymization cycle. To maximize the statistical effectiveness of our approach, we also presented two enhancements, namely embedding of complex business knowledge and runtime heuristics. We validated the approach on real-world central bank data. As future work, we plan to further enhance the framework, and test it in a variety of other real-world scenarios.

REFERENCES

- [1] Serge Abiteboul, Richard Hull, and Victor Vianu. 1995. *Foundations of Databases*. Addison-Wesley.
- [2] Paolo Atzeni, Luigi Bellomarini, Michela Iezzi, Emanuel Sallinger, and Adriano Vlad. 2020. Weaving Enterprise Knowledge Graphs: The Case of Company Ownership Graphs. In *EDBT*. 555–566.
- [3] Paolo Atzeni, Luigi Bellomarini, Paolo Papotti, and Riccardo Torlone. 2019. Meta-mappings for schema mapping reuse. *PVLDB* (2019).
- [4] Brett K. Beaulieu-Jones, Zhiwei Steven Wu, Chris Williams, Ran Lee, Sanjeev P. Bhavnani, James Brian Byrd, and Casey S. Greene. 2019. Privacy-Preserving Generative Deep Neural Networks Support Clinical Data Sharing. *Circ Cardiovasc Qual Outcomes* 12, 7 (07 2019).
- [5] Luigi Bellomarini, Davide Benedetto, Georg Gottlob, and Emanuel Sallinger. 2020. Vadalogue: A modern architecture for automated reasoning with large knowledge graphs. *Inf. Syst.* (2020), 101528.
- [6] Luigi Bellomarini, Emanuel Sallinger, and Georg Gottlob. 2018. The Vadalogue System: Datalog-based Reasoning for Knowledge Graphs. *PVLDB* 11, 9 (2018), 975–987.
- [7] Roberto Benedetti and Luisa Franconi. 1998. Statistical and technological solutions for controlled data dissemination. In *Pre-proceedings of New Techniques and Technologies for Statistics*, Vol. 1. 225–232.
- [8] Michael Benedikt, Bernardo Cuenca Grau, and Egor V Kostylev. 2018. Logical foundations of information disclosure in ontology-based data integration. *Artificial Intelligence* 262 (2018), 52–95.
- [9] Thijs Benschop, Cathrine Machingaut, and Matthew Welch. 2019. Statistical Disclosure Control: A Practice Guide. (2019).
- [10] Angela Bonifati, Ugo Comignani, and Efthymia Tsamoura. 2019. MapRepair: Mapping and Repairing under Policy Views. In *Proceedings of the 2019 International Conference on Management of Data*. 1873–1876.
- [11] Andrea Cali, Georg Gottlob, and Michael Kifer. 2013. Taming the Infinite Chase: Query Answering under Expressive Relational Constraints. *JAIR* 48 (2013), 115–174.
- [12] Andrea Cali, Georg Gottlob, Thomas Lukasiewicz, and Andreas Pieris. 2011. Datalog+/-: A Family of Languages for Ontology Querying. In *Datalog Reloaded*.
- [13] Peter Christen. 2012. *Data Matching - Concepts and Techniques for Record Linkage, Entity Resolution, and Duplicate Detection*. Springer.
- [14] Margareta Ciglic, Johann Eder, and Christian Koncilia. 2014. k-anonymity of microdata with NULL values. In *International Conference on Database and Expert Systems Applications*. Springer, 328–342.
- [15] Rémy Delanaux, Angela Bonifati, Marie-Christine Rousset, and Romuald Thion. 2018. Query-based linked data anonymization. In *International Semantic Web Conference*. Springer, 530–546.
- [16] Jörg Drechsler and Jerome Reiter. 2011. An empirical evaluation of easily implemented, nonparametric methods for generating synthetic datasets. *Computational Statistics & Data Analysis* 55 (2011), 3232–3243.
- [17] Cynthia Dwork. 2006. Differential Privacy. In *Automata, Languages and Programming*, Michele Bugliesi, Bart Preneel, Vladimiro Sassone, and Ingo Wegener (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 1–12.
- [18] Mark Elliot and Josep Domingo-Ferrer. 2018. The future of statistical disclosure control. *The National Statistician's Quality Review* (2018).
- [19] Mark J Elliot, Anna M Manning, and Rupert W Ford. 2002. A computational algorithm for handling the special uniques problem. *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems* 10, 05 (2002), 493–509.
- [20] Ronald Fagin, Phokion G. Kolaitis, Renée J. Miller, and Lucian Popa. 2003. Data Exchange: Semantics and Query Answering. In *ICDT*.
- [21] Ronald Fagin, Phokion G. Kolaitis, Renée J. Miller, and Lucian Popa. 2005. Data exchange: Semantics and query answering. *TCS* 336, 1 (2005), 89–124.
- [22] Luisa Franconi and Silvia Polettini. 2004. Individual risk estimation in μ -Argus: A review. In *Int. Workshop on Privacy in Statistical Databases*. 262–272.
- [23] Georg Gottlob, Thomas Lukasiewicz, and Andreas Pieris. 2014. Datalog+/-: Questions and Answers. In *KR*.
- [24] Bernardo Cuenca Grau and Egor V Kostylev. 2019. Logical foundations of linked data anonymisation. *J. of AI Research* 64 (2019), 253–314.
- [25] Paolo Guagliardo and Leonid Libkin. 2019. On the Codd semantics of SQL nulls. *Inf. Syst.* 86 (2019), 46–60.
- [26] Anco Hundepool, Josep Domingo-Ferrer, Luisa Franconi, Sarah Giessing, Eric Schulte Nordholt, Keith Spicer, and Peter-Paul De Wolf. 2012. *Statistical disclosure control*. John Wiley & Sons.
- [27] Anco Hundepool, A Van de Wetering, Ramya Ramaswamy, Peter-Paul de Wolf, Sarah Giessing, Matteo Fischetti, Juan-José Salazar, Jordi Castro, and Philip Lowthian. 2005. *r-argus user's manual*, version 3.3. *Statistics Netherlands, Voorburg, The Netherlands* (2005).
- [28] Michela Iezzi. 2020. Practical Privacy-Preserving Data Science With Homomorphic Encryption: An Overview. *CoRR* abs/2011.06820 (2020).
- [29] Anna M. Manning, David J. Haglin, and John A. Keane. 2008. A recursive search algorithm for statistical disclosure assessment. *Data Mining and Knowledge Discovery* 16, 2 (01 Apr 2008), 165–196.
- [30] Gregory Matthews and Ofer Harel. 2011. Data confidentiality: A review of methods for statistical disclosure limitation and methods for assessing privacy. *Stat. Surv.* 5 (01 2011).
- [31] Alan Nash and Alin Deutsch. 2007. Privacy in GLAV information integration. In *International Conference on Database Theory*. Springer, 89–103.
- [32] Aleksander Øhrn and Lucila Ohno-Machado. 1999. Using Boolean reasoning to anonymize databases. *Artificial intelligence in medicine* 15 3 (1999), 235–54.
- [33] Fabian Prasser and Florian Kohlmayer. 2015. Putting statistical disclosure control into practice: The ARX data anonymization tool. In *Medical Data Privacy Handbook*. Springer, 111–148.
- [34] Pierangela Samarati and Latanya Sweeney. 1998. Protecting privacy when disclosing information: k-anonymity and its enforcement through generalization and suppression. (1998).
- [35] Chris Skinner, Catherine Marsh, Stan Openshaw, and Colin Wymer. 1994. Disclosure control for census microdata. *JOS* 10, 1 (1994), 31–51.
- [36] Latanya Sweeney. 1997. Guaranteeing anonymity when sharing medical data, the Datafly System. *Proc. AMIA Fall Symposium* (1997), 51–55.
- [37] Latanya Sweeney. 2002. Achieving k-anonymity privacy protection using generalization and suppression. *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems* 10, 05 (2002), 571–588.
- [38] Matthias Templ, Alexander Kowarik, and Bernhard Meindl. 2013. *sdcMicro: Statistical Disclosure Control methods for the generation of public-and scientific-use files. Manual and Package* (2013).